

# **Strategies for Handling Missing Data**

Stephen Haptonstahl  
Berico Technologies

2012/08/30

# The Problem

# Data, in theory

Suppose we want to mine sales records to see what kind of customer buys more from us.

<b>sales</b>	<b>female</b>	<b>education</b>	<b>income</b>
1209	0	high school	35000
2587	1	4yr degree	45000
13265	1	grad degree	65000
2298	0	some college	49000

# Data, in practice

Suppose we want to mine sales records to see what kind of customer buys more from us.

<b>sales</b>	<b>female</b>	<b>education</b>	<b>income</b>
1209	0	high school	35000
2587	NA	4yr degree	45000
13265	1	NA	65000
NA	0	some college	49000

# Who cares?

- Best case if you ignore missing data: throwing away data
- Same as loss of efficiency
- Less certainty (less statistical significance)
- **Some patterns will be completely missed**
- Data scientists are replaced by equally incompetent algorithms, ie.
- Mass hysteria

# Today's talk

- The Problem
- Typical Method Versus a Better Method
- Naive Imputation
- Home-Rolled Imputation
- Existing Tools
  - *Small Data Sets*
  - *Large Data Sets*
  - *Real-time*
- When Imputation Fails
- Fin

# Typical Method Versus a Better Method

# (The Evil that is) Row Deletion

- Idea: Only include full rows in the analysis.
- This is the most common method.
  - *Default for most software, including `lm` and other `R` functions.*
- Throws away data, perhaps a lot.
  - *Ex: 50 fields for each customer*
  - *How many have all 50 observed? (Ans: none except test data)*
- **Some patterns will be completely missed**
- We know where this leads



# What is imputation?

## Imputation

Filling in missing data with guesses, good or bad, for what would have been observed.

- Example 1: “These scientists all say the earth is warming, but today it’s cold. They must be in a conspiracy!”
- Example 2: “These scientists all say the earth is warming, but today it’s cold. Perhaps they only measured on warm days.”
- Both are wrong, but one is a better (more likely to be true) guess than the other.
- The better guess will be less likely to lead you to bad decisions, algorithms, hysteria, etc.

# When is imputation bad

*If only we knew what was there. Then we would be able to make a good guess at what was there.*

- Missingness process **not ignorable**, which means
- Missingness depends on unobserved data
- Types of data-generating processes
  1. *Missing Completely At Random (MCAR): no data will help you predict which values are missing*
    - Random noise
  2. *Missing at Random (MAR): observed values can help you predict which values are missing*
    - Faulty switches
  3. *Not ignorable (NMAR): unobserved values can help you predict which values are missing*
    - Losses reported as zero income (censored)
- Imputing when missingness is not ignorable leads to bias, bad decisions, etc.
- cf. Donald Rubin (1976)

# Naive Imputation

# Mean, Median, Mode

*We just filled in the missing values with zero.*

How we guess (hopefully) depends on what we expect to be there.

- Continuous variable

- *floating point*
- *Use mean*

- Ordinal (disagree, neither agree nor disagree, agree): median

- Count (or other integer): median

- Categorical (states; colors): mode

```
StatMode <- function(x) names(table(x))  
[which.max(table(x))]
```

# Mean, Median, Mode (continued)

Suppose we want to mine sales records to see what kind of customer buys more from us.

<b>sales</b>	<b>female</b>	<b>education</b>	<b>income</b>
1209	0	high school	35000
2587	NA	4yr degree	45000
13265	1	NA	65000
NA	0	some college	49000

# Mean, Median, Mode (continued)

Suppose we want to mine sales records to see what kind of customer buys more from us.

<b>sales</b>	<b>female</b>	<b>education</b>	<b>income</b>
1209	0	high school	35000
2587	<b>0</b>	4yr degree	45000
13265	1	some college	65000
<b>5687</b>	0	<b>some college</b>	49000
<b>mean</b>	<b>mode</b>	<b>median</b>	

# Mean, Median, Mode (continued)

## ■ Pros

- *Fast*
- *Principled*
- *Better than throwing away data*
- *Tends to work*

## ■ Cons

- *We can do better.*
- *“Making up data” makes us more certain than we should be.*

# Home-Rolled Imputation



# Imputation with Regression

1. Identify a cell  $X[i, j]$  to impute.
2. Find all rows `prediction.rows` that:
  1. *have an observed value for column  $j$*
  2. *have observed values for all other columns `obs.cols` that row  $i$  has*
3. Using those rows: `reg <- lm(X[prediction.rows, j] ~ X[prediction.rows, obs.cols])`
4. Fill with `predict(reg, names(X)[obs.cols]=X[i, obs.cols])`

# Imputation with Regression (continued)

<b>y</b>	<b>x1</b>	<b>x2</b>
7.44	1.21	1.95
3.25	-1.35	0.11
3.90	NA	0.63
5.19	0.16	2.27
5.57	0.17	1.53
...	...	...

# Imputation with Regression (continued)

<b>y</b>	<b>x1</b>	<b>x2</b>
7.44	1.21	1.95
3.25	-1.35	0.11
3.90	-0.89	0.63
5.19	0.16	2.27
5.57	0.17	1.53
...	...	...

- Missing value: -0.89

- Predicted value: -0.89

# Type of Regression Varies by Variable Type

## Continuous

Linear regression with `lm( )`

## Ordinal

Ordered logit/probit regression with `MASS::polr`

## Count

Poisson regression with `glm(..., family=poisson(link = "log"))`

## Categorical

Multinomial logit with `nnet::multinom`

# Safety First!

- Look out for having too little data to estimate. (Default to mean/median/mode?)
- Data is dirty in other ways (illegal values).
- Iterative algorithms sometimes don't converge.
- Still have problem that “Making up data” makes us more certain than we should be.

# Existing Tools

# Small Data Sets

`Amelia` (King) and `mi` (Gelman) work differently from home-rolled solutions.

## Multiple Imputation

1. Generate  $m$  copies of the data set, each with **different** imputed values.
2. Perform the desired analysis on each imputed data set to get the quantity of interest  $q$ .
3. Aggregate the results:
  1. *If  $m = 5$  or so*
    - Point-estimates: `mean(q)`
    - SEs: fairly simple formula
  2. *If  $m$  large*
    - Treat as simulation result
    - `mean` and `sd`
4. Check distribution of observed versus imputed values (`plot`, `ks.test`)

# Amelia Example

```
library(Amelia)
data(africa)
```

```
a.out <- amelia(x = africa, cs = "country", ts = "year", logs = "gdp_pc")
```

```
## -- Imputation 1 --
##
##  1  2  3
##
## -- Imputation 2 --
##
##  1  2  3
##
## -- Imputation 3 --
##
##  1  2  3
##
## -- Imputation 4 --
##
##  1  2
##
## -- Imputation 5 --
##
##  1  2
##
```



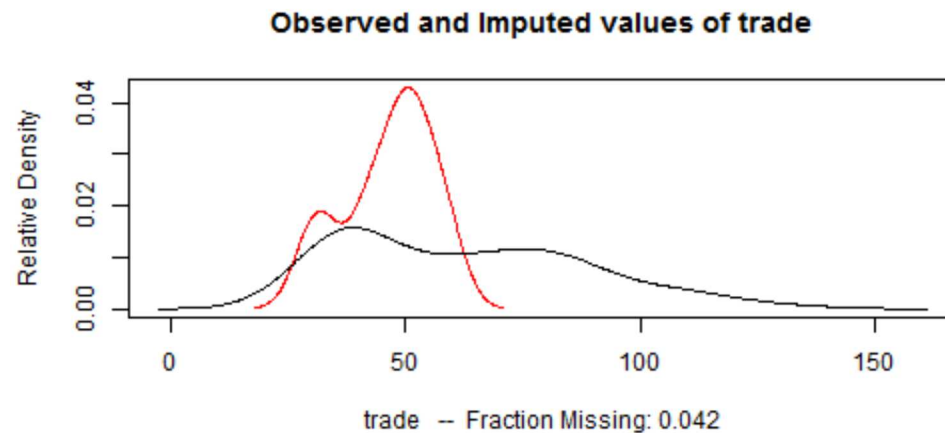

# Amelia Example (continued)

```
summary(a.out)
```

```
##
## Amelia output with 5 imputed datasets.
## Return code: 1
## Message: Normal EM convergence.
##
## Chain Lengths:
## -----
## Imputation 1: 3
## Imputation 2: 3
## Imputation 3: 3
## Imputation 4: 2
## Imputation 5: 2
##
## Rows after Listwise Deletion: 115
## Rows after Imputation: 120
## Patterns of missingness in the data: 3
##
## Fraction Missing for original variables:
## -----
##
##          Fraction Missing
## year                0.00000
## country              0.00000
## gdp_pc               0.01667
## infl                 0.00000
## trade                0.04167
## civlib               0.00000
## population           0.00000
##
```

# Amelia Example (continued)

```
plot(a.out)
```



# That's hard. What have we gained or lost?

- Correct reporting of uncertainty (SEs)
- Actually, it's easier than rolling your own
  - *Smart about pathologies*
  - *Smart about variable types*
  - *Code is shorter if you are implementing in R*
  - *Very fast on reasonably sized data sets*
- If you are **not** implementing in R, you have a lot more work
  - *Reading the paper(s), implementing your own version*
  - *Quite do-able, not for faint at heart*
- Not scalable

# Large Data Sets: Bayes

## Bayesian data augmentation

- When using Bayes to estimate a model, put a prior on the *data*.
- The algorithm will give draws from the posterior distribution of the missing values.
- It uses the model defined (whatever it is) to impute.

# Large Data Sets: Bayes (continued)

## Problems

- Slow
- Convergence can be tricky
- Expertise required for custom samplers, but libraries for implementing custom solutions: `rcppbugs` in R, others for C++, Python, etc.

# Large Data Sets: Bayes (continued)

Example: Using Bayes for IRT on network links to generate DILS (data-informed link strength)

linkid	net1	net2	net3	...	netk
1	1	0	0	...	1
2	0	0	1	...	0
3	1	NA	1	...	1
4	1	0	NA	...	0
5	NA	0	0	...	1

becomes...

# Large Data Sets: Bayes (continued)

Example: Using Bayes for IRT on network links to generate DILS (data-informed link strength)

linkid	dils
--------	------

1	.12
---	-----

2	.01
---	-----

3	.34
---	-----

4	.08
---	-----

5	.03
---	-----

# Large Data Sets: Random Forest Classifier

Easy solution: use `missForest`

```
library(Amelia) # to provide the data
data(africa)
library(missForest)
```

```
africa.rf <- missForest(africa)
```

```
## missForest iteration 1 in progress...done!
## missForest iteration 2 in progress...done!
## missForest iteration 3 in progress...done!
## missForest iteration 4 in progress...done!
```

- Yields very good point predictions (single imputation).
- Scales well.
- Requires a **lot** of data (think non-parametric).
- Without care might give silly values.
- Still should check imputed distributions against observed distributions.



# Large Data Sets: Random Forest Classifier (continued)

- If you want good measures of uncertainty (multiple imputation) then you'll have to find some way to keep the `ntrees` individual trees used in each of the `nvars` forests.
- Run `randomForest(..., keep.forest=TRUE)` with each column as dependent variable.
  - Result: *`ntrees` × `nvars` trees*
- Use `lapply` to get `ntrees` imputed data sets.
  - Create *`ntrees` target copies of the original data set.*
  - For each column use the forest which had that column as the dependent variable.
  - For each empty cell generate a prediction for each tree in that forest.
  - Store the result for each tree in the corresponding target copy.
- Run your analysis over each of the imputed data sets.
- Aggregate the results using `mean` and `sd`.

# Real-Time: FastImputation

- Project: Improve fraud detection classifier
- Same model as Amelia
- Learns (slowly) using lots of data
- Imputes a single row at a time (very, well, fast)
- Available on CRAN
- Caveat emptor (as with all R packages)

# When Imputation Fails

# To Impute or Not?

## ■ Assumptions of Imputation

- *Data is MCAR or MAR+PD*
- *“Small” fraction of data missing*
  - *< 10%: usually no problem*
  - *10%-20%: Be careful*
  - *20%: There be dragons*

## ■ Always\* better than row deletion

# What to Do When Imputation Fails

- Explicitly model the reason observations are missing

- *Tobit (incomes censored, displayed as zero)*
- *Truncated regression (heights in military)*
- *Abstentions in voting (more likely if you disagree with party)*

or

- Impute anyway and warn the consumer (worth considering)

# Fin

# Links and Contact Info

- This presentation and code: [https://github.com/shaptonstahl/MissingData\\_2012-08](https://github.com/shaptonstahl/MissingData_2012-08)
- Amelia: <http://gking.harvard.edu/amelia> (the 2001 paper is a good place to start)
- mi: <http://cran.r-project.org/web/packages/mi>
- Great ref on missing data, imputation: <http://gking.harvard.edu/files/evil.pdf>
- FastImputation: <http://cran.r-project.org/web/packages/FastImputation>
- Data-Informed Link Strength (DILS): <https://github.com/shaptonstahl/dils>
- Great place to work: <http://www.bericotechnologies.com>
- Stephen Haptonstahl – [srh@haptonstahl.org](mailto:srh@haptonstahl.org) – @polimath